

# RuHuman: A Resilient Multimodal AI-Powered Audio Verification System

Timothy Do

Friday, December 15th, 2023, 7:00 PM PST

ECE202A Embedded Systems @ UCLA



# Motivation and Objectives

- Audio deep fakes are becoming more easily spoofable with the advent of generative AI.
  - Growing number in deep fake crimes
- RuHuman evaluates existing audio liveness detection systems that exploit multiple audio modalities.
- The Objectives of this RuHuman Are:
  - Evaluating different audio liveness detector architectures against state of the art voice cloners.
  - Developing a user interface to make gateway devices (e.g. phone) easily verify with their microphones



# Previous/Related Work

- ASVSpooF is a dedicated challenge that to develop audio spoofing classifiers. Three main challenges are PA, LA, and DF (RuHuman focuses on LA).
- Researchers take advantage of various encodings (STFT, MFCC, CQCC) to augment the training resolution of the audio sample.
- UCLA NESL developed ResNet models with different embeddings and performed model fusion to minimize ASV/CM loss metrics in ASVSpooF 2019.




Figure 1: Model architecture for the Spec-ResNet model. Detailed structure of residual blocks is shown in 2.

# Novelty of Project

- Datasets used for training in ASVSpooof2019 LA challenge were in a controlled environment, insusceptible to 'out in the wild' samples
- RuHuman investigates the performance of different detector architectures when additive noise is mixed in with the suspected audio sample.
- Additionally, RuHuman implements an interactive user interface that allow for a wide range of computing (from a cell phone to a computer workstation) to access multiple detector models with just an HTTP Post Request.

## RuHuman Audio Verification System

Check if Your Audio is Produced by an AI:

Option 1: Record a Live Sample

Record

Option 2: Upload a Recording

Choose File No file chosen

Verify Audio Reset

Choose a Detector Model: Tortoise



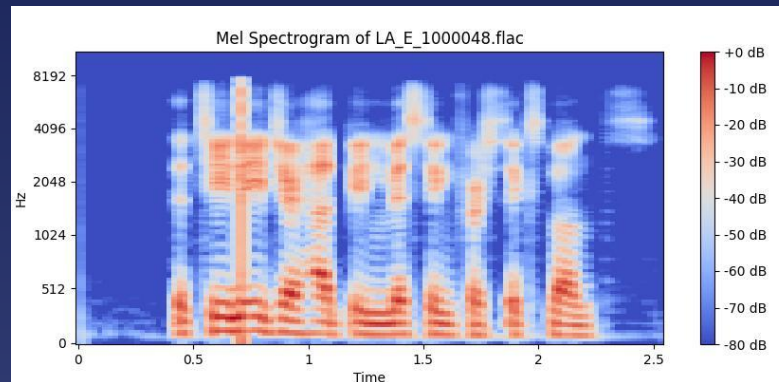
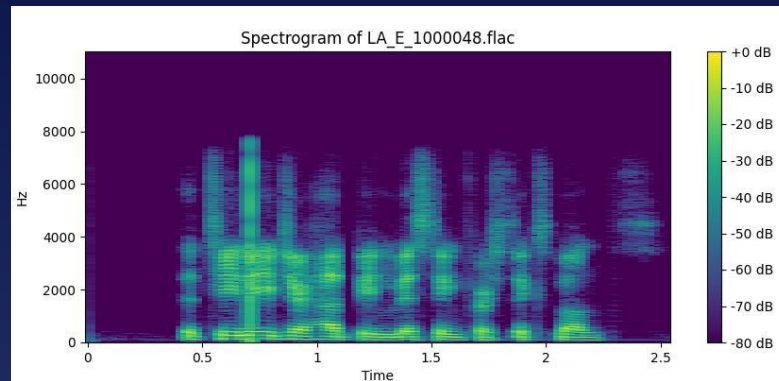
# Background: Audio Feature Extractions

**STFT** (Short-Time Fourier Transform): Compute FFT over a window, use log to plot Spectrogram.

**LFCC** (Linear-Frequency Cepstral Coefficients): Map frequency powers through a linear filter bank then take DCT of log of filter coefficients

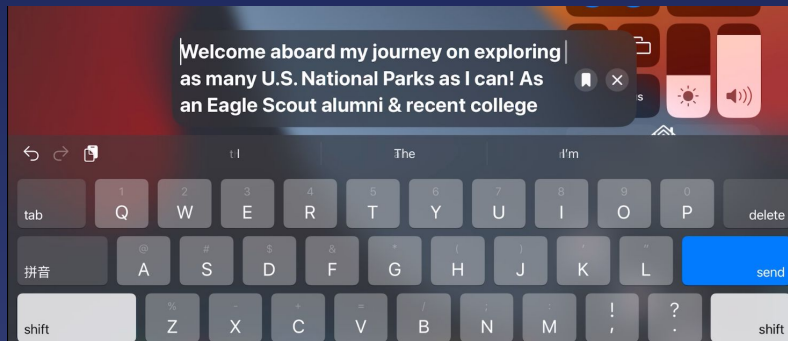
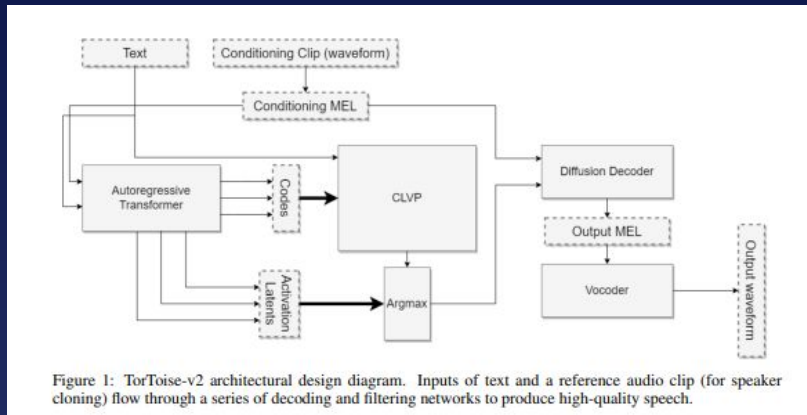
**MFCC** (Mel-Frequency Cepstral Coefficients): Map frequency powers through a mel filter bank (tuned for humans), then take DCT of coefficients.

**CQCC** (Constant Q Cepstral Coefficients): Uses CQT basis for cepstrum. More spectro/temporal resolution at low/high frequencies.



# Overview of Voice Cloners

- **Tortoise:** A voice cloner using Autoregressive Diffusive decoder with DDPM noise schedulers.
  - Pre-trained on 49,000 hours of audio clips from podcasts/videos.
- **Apple Personal Voice:** A proprietary personalized voice synthesizer that is trained by using 15 mins of live utterance
  - Trained locally on iOS device, use text to speech to synthesize audio.
  - Aside: Not easily exportable, must use screen recorder to grab audio.



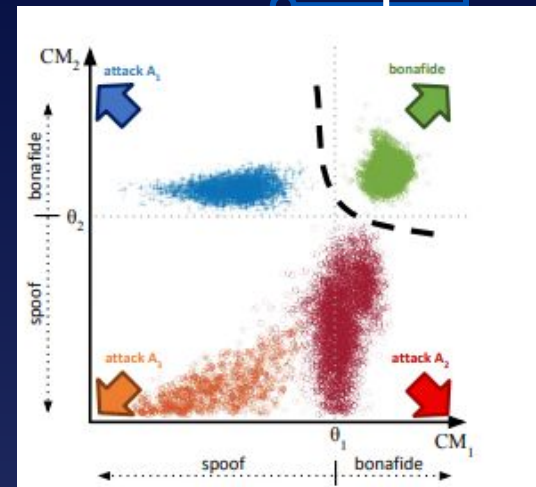
# Overview of Voice Detectors

## Baseline ASVSpoof Models:

- **LFCC-GMM:** A Two-Class Gaussian Mixture Model with LFCC input. 70-Channel Linear Filters. Outputs LLR between bonafide & spoof.
- **CQCC-GMM:** A Two-Class Gaussian Mixture Model with CQCC input. Features are 1st 19 Coeff. Outputs LLR between bonafide & spoof.
- **RawNet2:** An End to End Neural Network using mel sinc filters. Outputs two-class likelihood between bonafide and spoof.
- All Baseline models trained with ASVSpoof2019 Training Set (~25380 Samples, 20/107 Distinct Speakers from total set)

## In the Wild Models:

- **Tortoise:** An Audio Mini Encoder Classifier. Outputs probability on whether an audio sample was generated by TortoiseTTS.
  - Trained on same data used by detector (49,000 hours of podcasts/videos from internet).



Layer	Input: 64000 samples	Output shape
Fixed Sinc filters	Conv(129,1,128)	(21290,128)
	Maxpooling(3)	
	BN & LeakyReLU	
Res block	BN & LeakyReLU Conv(3,1,128) BN & LeakyReLU Conv(3,1,128) Maxpooling(3) FMS	× 2 (2365,128)
Res block	BN & LeakyReLU Conv(3,1, 512) BN & LeakyReLU Conv(3,1, 512) Maxpooling(3) FMS	× 4 (29,512)
GRU	GRU(1024)	(1024)
FC	1024	(1024)
Output	1024	2



# Overview of User Interface

- Web Server Implemented in Python Flask to serve many gateway devices that have internet access.
- Users can verify their audio sample in 2 ways:
  - **1.** Users record a live sample client-side by using `getUserMedia()` from Javascript, which then is saved to a blob.
  - **2.** Users upload a pre-recorded audio file from their device to the HTML form to post.
- Users upload data via HTTP request post to server.
- File is stored and pass through detector via refactored callable Python functions (GMMs: `matlab.engine`)
- Dashboard displays uploaded audio (verifying it's on the server), in addition to visualizing it's encodings with *librosa* and outputting metrics of verification.

## RuHuman Audio Verification System

Check if Your Audio is Produced by an AI:

Option 1: Record a Live Sample

Record

Option 2: Upload a Recording

Choose File No file chosen

Verify Audio Reset

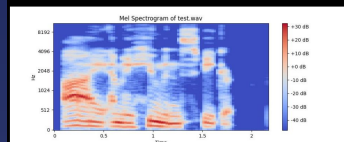
Choose a Detector Model: Tortoise

Your Uploaded Audio Sample: (test.wav)

0:00:00



Features:



Detector Model: Tortoise (Detected in 0.16 seconds)

Probability of Being Spoofed by Generative AI via Tortoise Model: 99.98%

[Return to Homepage](#)



# Metrics of Evaluation

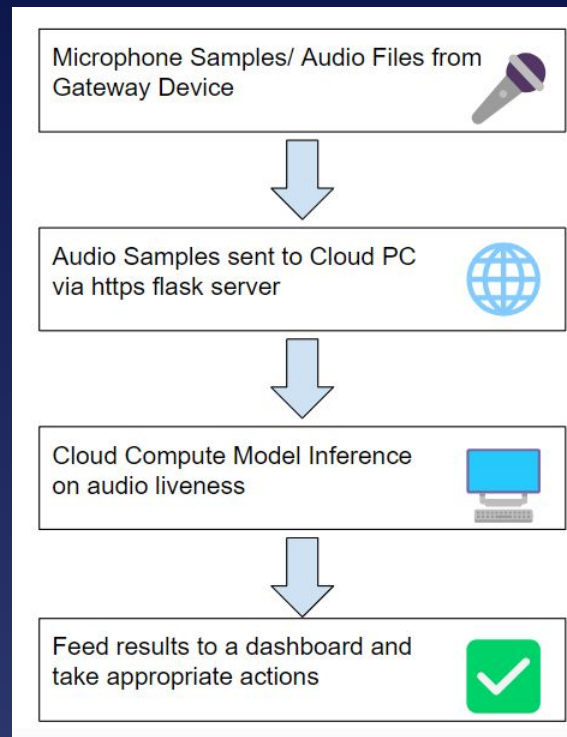
- **t-DCF:** A custom cost function that weighs between predefined ASV scores (from organizers), speaker priors, and CM scores provided by the detector architecture.
- **EER:** The likelihood threshold at which the miss rate and the false alarm rate equate.
- **Computation Time:** The duration for a single audio sample to be verified.
  - Ideally in a live setting should be less than duration of the sample.
- Objective is to minimize all metrics for accuracy and efficiency in live settings.

$$\begin{aligned} \text{t-DCF}(s, t) = & C_{\text{miss}}^{\text{asv}} \cdot \pi_{\text{tar}} \cdot P_a(s, t) \\ & + C_{\text{fa}}^{\text{asv}} \cdot \pi_{\text{non}} \cdot P_b(s, t) \\ & + C_{\text{fa}}^{\text{cm}} \cdot \pi_{\text{spooof}} \cdot P_c(s, t) \\ & + C_{\text{miss}}^{\text{cm}} \cdot \pi_{\text{tar}} \cdot P_d(s). \end{aligned}$$

- $C_{\text{miss}}^{\text{asv}}$  – cost of ASV system rejecting a target trial.
- $C_{\text{fa}}^{\text{asv}}$  – cost of ASV system accepting a nontarget trial.
- $C_{\text{miss}}^{\text{cm}}$  – cost of CM rejecting a human trial.
- $C_{\text{fa}}^{\text{cm}}$  – cost of CM accepting a spoof trial.

# Technical Approach

- **Phase 1: Investigation of Additive Noise Sources** (ASVSpooof Submissions Only)
  - Work with each baseline detector to produce metrics on the ASVSpooof2021 Evaluation Set (181,566 samples, 67 Distinct Speakers, 37 F, 30 M).
  - Inject each audio sample in the evaluation set with additive white gaussian noise (SNR = 50).
  - Recompute metrics with infected evaluation set and compare scores.
- **Phase 2: Deployment of User Interface**
  - Send/receive audio from client/gateway with multiple upload options.
  - Refactor detectors to be callable functions with audio path argument.
  - Show audio features and results on dashboard.



# Primary Metric Metric: t-DCF & EERs

Detector Architecture (Normal)	Pooled t-DCF	Pooled EER (%)
Baseline LFCC-GMM (MATLAB)	0.5758	19.30
Baseline CQCC-GMM (MATLAB)	0.4964	15.62
Baseline RawNet2 (Python)	0.4255	9.49

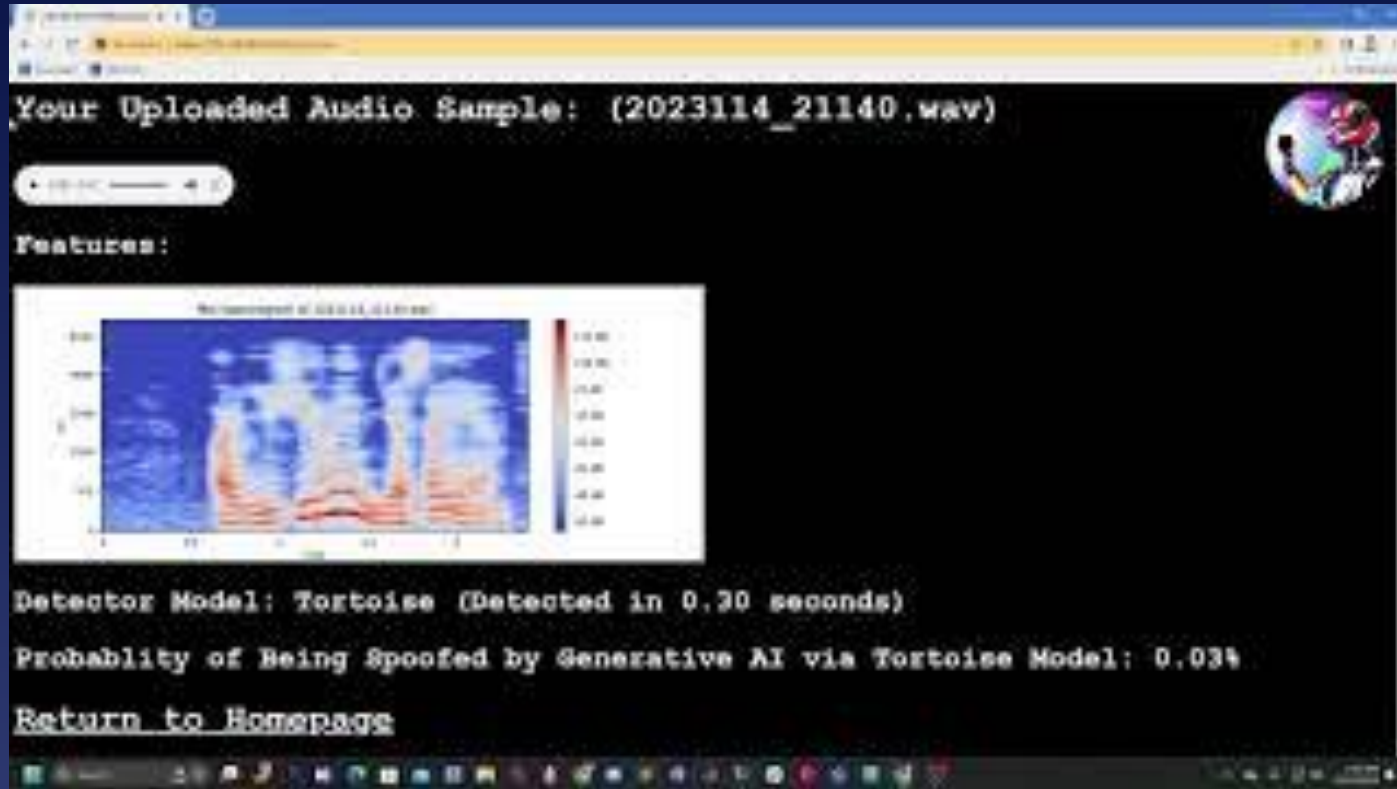
Detector Architecture (AWGN)	Pooled t-DCF	Pooled EER (%)
Baseline LFCC-GMM (MATLAB)	0.8527	41.33
Baseline CQCC-GMM (MATLAB)	0.7609	31.99
Baseline RawNet2 (Python)	0.3317	7.23

# Secondary Metric: Computation Time

Detector Architecture	Average Computation Time per Sample (s)
Tortoise Audio Mini Encoder (Python)	0.22
Baseline LFCC-GMM (MATLAB)	0.18
Baseline CQCC-GMM (MATLAB)	0.27
Baseline RawNet2 (Python)	0.30

*\*on average each sample was approximately 3-5 seconds long*

# Demo



# Conclusion & Future Directions

## Conclusion:

- Primary evaluation from phase 1 indicates that additive noise can be injected to GMMs enhance a plausible spoofing attack to being successful.
  - Metrics actually went lower for RawNet2, perhaps the initial layer of Mel Sinc Filters mitigated the effect of AWGN.
- All detectors were able to detect an audio samples within a second, making them usable for liveness detection.
- User Interface was deployed and easily accessible from gateway devices.

## Future Directions:

- Preprocess audio samples with denoising to isolate speaker from noise before being sent to spoof detection model.
- Perform embedded actions on classifier outputs (e.g. turn a light red/green)
- Audio spoofing detection can be fused with other modalities of liveness detection (i.e. video deep fake detection) for a complete system. May experiment with this pipeline over winter break on a separate branch.

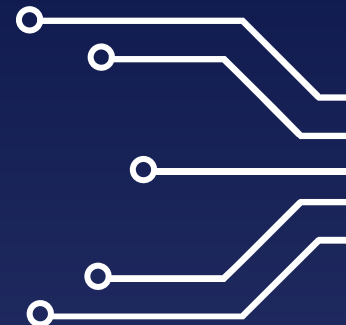
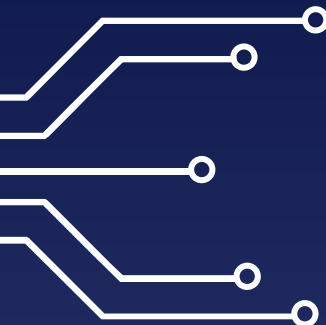
# References

- [1] X. Liu et al., "ASVspoof 2021: Towards Spoofed and Deepfake Speech Detection in the Wild," in IEEE/ACM Transactions on Audio, Speech, and Language Processing, vol. 31, pp. 2507-2522, 2023, doi: 10.1109/TASLP.2023.3285283.
- [2] M. Alzantot, Z. Wang, and M. B. Srivastava, "Deep Residual Neural Networks For Audio Spoofing Detection," 2019, arXiv:1907.00501.
- [3] Yamagishi, Junichi; Todisco, Massimiliano; Sahidullah, Md; Delgado, Héctor; Wang, Xin; Evans, Nicolas; Kinnunen, Tomi; Lee, Kong Aik; Vestman, Ville; Nautsch, Andreas. (2019). ASVspoof 2019: The 3rd Automatic Speaker Verification Spoofing and Countermeasures Challenge database, [sound]. University of Edinburgh. The Centre for Speech Technology Research (CSTR). <https://doi.org/10.7488/ds/2555>.
- [4] A. Paszke et al., "PyTorch: An Imperative Style, High-Performance Deep Learning Library", in Proceedings of the 33rd International Conference on Neural Information Processing Systems, Red Hook, NY, USA: Curran Associates Inc., 2019.
- [5] J. Betker, 'Better speech synthesis through scaling', arXiv [cs.SD]. 2023.



# References

- [6] T. Kinnunen et al., "t-DCF: a Detection Cost Function for the Tandem Assessment of Spoofing Countermeasures and Automatic Speaker Verification", arXiv [eess.AS]. 2019.
- [7] "Apple previews Live Speech, Personal Voice, and more new accessibility features," Apple Newsroom, <https://www.apple.com/newsroom/2023/05/apple-previews-live-speech-personal-voice-and-more-new-accessibility-features/>.
- [8] M. Todisco, H. Delgado and N. Evans, "Constant Q cepstral coefficients: a spoofing countermeasure for automatic speaker verification", Computer, Speech and Language, vol. 45, pp. 516 –535, 2017.
- [9] H. Tak, J. Patino, A. Nautsch, N. Evans, M. Todisco, "Spoofing Attack Detection using the Non-linear Fusion of Sub-band Classifiers" in Proc INTERSPEECH, 2020.
- [10] H. Tak, J. Patino, M. Todisco, A. Nautsch, N. Evans, and A. Larcher, 'End-to-End anti-spoofing with RawNet2', in IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), 2021, pp. 6369–6373.



**Thanks for Listening!**  
**Any Questions?**

